

Clamp - Seamless integration of Python and Java

Jim Baker, Rackspace

jim.baker@rackspace.com

Clamp background

Clamp -
Seamless
integration of
Python and
Java

Jim Baker,
Rackspace

- Part of the jythontools project
(<https://github.com/jythontools>)
- Improve Python \Leftrightarrow Java integration (which is already very good)
- Enable precise layout of the generated Java bytecode for Python classes
- So they can be used as modern Java classes - annotation metadata, type signatures
- Jar packaging into site-packages
- Or entire Jython installation wrapped into a **single jar**

Benefits

Clamp -
Seamless
integration of
Python and
Java

Jim Baker,
Rackspace

- JVM frameworks can readily work with clamped code, oblivious of its source
- Especially need single jar support
- Developers can stay as much in Python as possible
- Working on SQLAlchemy-like DSL

Development

Clamp -
Seamless
integration of
Python and
Java

Jim Baker,
Rackspace

- Heavily uses support for metaprogramming in Python and Java
- If there's a metaprogramming facility, we seem to be either using it now or exploring it
- Useful outcome already: `merge type, java.lang.Class` so equivalent types for metaclass usage

Example: Storm

Clamp -
Seamless
integration of
Python and
Java

Jim Baker,
Rackspace

- “Real-time” complex event processing system
- Runs topology of storms, bolts to process events (“tuples”)
- Can support at-least-once, exactly-once semantics

What is needed for Storm

Clamp -
Seamless
integration of
Python and
Java

Jim Baker,
Rackspace

- No-arg constructors
- Specify `serialVersionUID`
- Single jar support
- Needs to be able to resolve class names to classes
(`Class.forName`)

Python class, extending Java interfaces

Clamp -
Seamless
integration of
Python and
Java

Jim Baker,
Rackspace

```
from java.io import Serializable
from java.util.concurrent import Callable

class BarClamp(Callable, Serializable):

    def call(self):
        return 42
```

How to use from Java?

Clamp -
Seamless
integration of
Python and
Java

Jim Baker,
Rackspace

- Can use as a Java callback
- Can use JSR-223 or embed Jython runtime into Java
- But cannot directly use as-is from Java
- Specifically no way to construct a new instance of the class
- Definitive Guide to Jython goes into detail: object factories, etc
- Good if you are already using JSR-223 etc, want to support scripting, etc
- Bad if you just want to use Python code in your framework as one component, instead of having to write the component in Java
- Note that solutions like Scala, Clojure, face similar issues: they need to describe how they should be exposed with a Java API
- Solution: Clamp!

Python class, clamped

Clamp -
Seamless
integration of
Python and
Java

Jim Baker,
Rackspace

Simply add the Clamp base, a metaclass:

```
from java.io import Serializable
from java.util.concurrent import Callable
from clamp import clamp_base
```

```
BarBase = clamp_base("bar")  # Java package prefix
```

```
class BarClamp(BarBase, Callable, Serializable):

    def call(self):
        return 42
```

Clamping your class

Clamp -
Seamless
integration of
Python and
Java

Jim Baker,
Rackspace

Key insight: ahead-of-time builds through setuptools:

```
import ez_setup
ez_setup.use_setuptools()

from setuptools import setup, find_packages

setup(
    name = "clamped",
    version = "0.1",
    packages = find_packages(),
    install_requires = ["clamp>=0.3"],
    clamp = ["clamped"],
)
```

Using from Java

Clamp -
Seamless
integration of
Python and
Java

Jim Baker,
Rackspace

Simply import clamped Python classes into Java code!

```
import bar.clamped.BarClamp;

public class UseClamped {
    public static void main(String[] args) {
        BarClamp barclamp = new BarClamp();
        try {
            System.out.println("BarClamp: " +
                               barclamp.call());
        } catch (Exception ex) {
            System.err.println("Exception: " + ex);
        }
    }
}
```

Java/Python rendezvous

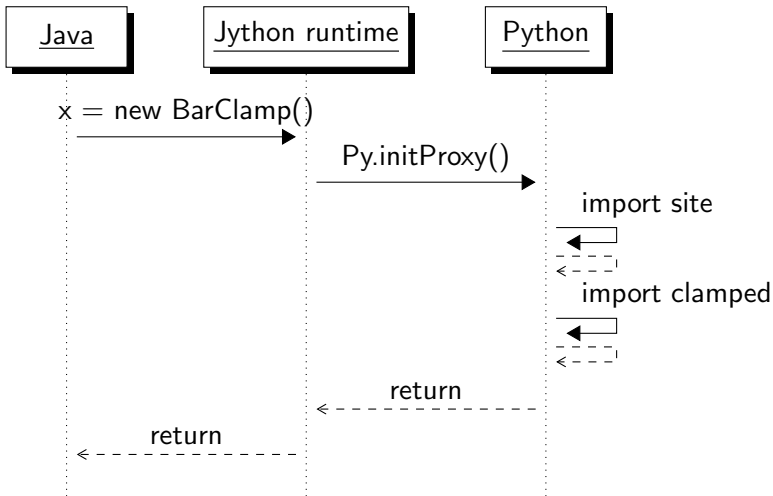
Clamp -
Seamless
integration of
Python and
Java

Jim Baker,
Rackspace

Initializing BarClamp

Clamp -
Seamless
integration of
Python and
Java

Jim Baker,
Rackspace



Loading Jython runtime

Clamp -
Seamless
integration of
Python and
Java

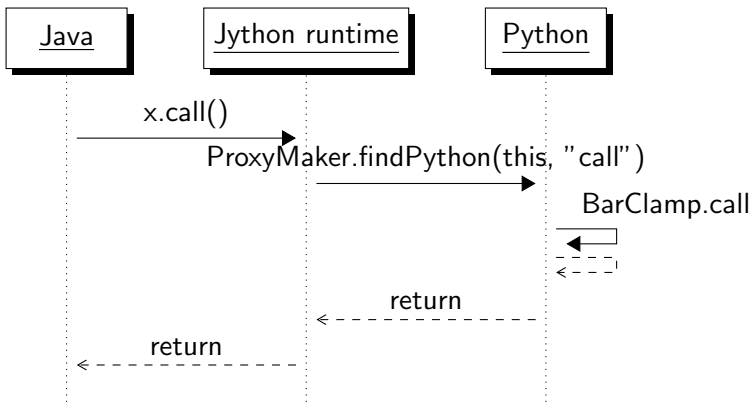
Jim Baker,
Rackspace

- When does the Jython runtime get loaded? (bootstrap problem!)
- `Py.initProxy` calls `Py.getThreadState`, which in turn loads the runtime as necessary!
- `ThreadState` extends `java.lang.ThreadLocal`

Making a call

Clamp -
Seamless
integration of
Python and
Java

Jim Baker,
Rackspace



Maintaining state

Clamp -
Seamless
integration of
Python and
Java

Jim Baker,
Rackspace

- ThreadState also maintains state through the call stack, from Java to Python and back, as necessary
- Other critical call: `findPython`, so that proxies can get the corresponding Python code!

Metaprogramming

Clamp -
Seamless
integration of
Python and
Java

Jim Baker,
Rackspace

- What do we mean by metaprogramming?
- Solution to every problem in CS is indirection. . .
- Except for indirection!

Better DSLs through metaclasses

Clamp -
Seamless
integration of
Python and
Java

Jim Baker,
Rackspace

- Declarative DSL
- Construct a mapper metaclass (similar to and inspired by SQLAlchemy)

API design aside

Clamp -
Seamless
integration of
Python and
Java

Jim Baker,
Rackspace

- Original clamp work was writing this in Java
- Getting too complex
- Let's use Python instead!
- Even if in places we are generating some Java bytecode!

Metaclass factory

Clamp -
Seamless
integration of
Python and
Java

Jim Baker,
Rackspace

Integrations can associate more info with metaclass base:

```
def clamp_base(package, proxy_maker=ClampProxyMaker):
    def _clamp_closure(package, proxy_maker):
        class ClampProxyMakerMeta(type):
            def __new__(cls, name, bases, d):
                d = dict(d)
                d['__proxymaker__'] = proxy_maker(
                    package=package)
                return type.__new__(cls, name, bases, d)
        return ClampProxyMakerMeta
    class ClampBase(object):
        __metaclass__ = _clamp_closure(
            package=package, proxy_maker=proxy_maker)
    return ClampBase
```

Proxy maker

Clamp -
Seamless
integration of
Python and
Java

Jim Baker,
Rackspace

- Heart of Clamp!
- Need to **precisely** generate in Java bytecode a Java class that implements interfaces and/or extends a class
- With a given name, because Java can be finicky about that (`Class.forName`, any static linkage)
- Constructors (no-arg at least)
- Static fields such as `serialVersionUID` to support `Serializable`
- Use the same bytecode!

__proxymaker__ protocol

Clamp -
Seamless
integration of
Python and
Java

Jim Baker,
Rackspace

- Specific to Jython
- New __proxymaker__ protocol allows for a CustomMaker to intercept the construction of a Java proxy
- Code generation
- Saving bytes
- Turning into a class via a ClassLoader

Code generation example: <clinit>

Clamp -
Seamless
integration of
Python and
Java

Jim Baker,
Rackspace

Every class has a specially-named hidden method, <clinit>, to support initialization for the class itself:

```
public static final long serialVersionUID;  
static {  
    serialVersionUID = 42L;  
}
```

- <clinit> automatically emitted by the Java compiler
- **Clamp needs to do the same**

Code generation, in Python

Clamp -
Seamless
integration of
Python and
Java

Jim Baker,
Rackspace

Assuming self.constants is initialized like so:

```
{ "serialVersionUID":  
  (java.lang.Long(42), java.lang.Long.TYPE), ... }
```


serialVersionUID and dynamic typing

Clamp -
Seamless
integration of
Python and
Java

Jim Baker,
Rackspace

- Q: What is the correct value of `serialVersionUID`?
- A: 1L - this constant is to support dynamic evolution, but dynamic typing does that for us anyway!

Code generation

Clamp -
Seamless
integration of
Python and
Java

Jim Baker,
Rackspace

Standard visitor approach:

```
def doConstants(self):
    code = self.classfile.addMethod("<clinit>",
        ProxyCodeHelpers.makeSig("V"), Modifier.STATIC)
    for constant, (value, constant_type) in sorted(
        self.constants.iteritems()):
        self.classfile.addField(
            constant,
            CodegenUtils.ci(constant_type),
            Modifier.PUBLIC | Modifier.STATIC |
            Modifier.FINAL)
        code.visitLdcInsn(value)
        code.putstatic(self.classfile.name,
            constant, CodegenUtils.ci(constant_type))
    code.return_()
```

Current DSL

Clamp -
Seamless
integration of
Python and
Java

Jim Baker,
Rackspace

<http://clamp.readthedocs.org>

Locating jars

Clamp -
Seamless
integration of
Python and
Java

Jim Baker,
Rackspace

- Works because of this importer in `sys.path`:
`'__classpath__'`
- Other `sys.path` magic allows for dynamic addition to the `sys.path`
- How can we do this?
- Jython uses custom `ClassLoader` objects - Java is very flexible with respect to how to find, load classes

pth support

Clamp -
Seamless
integration of
Python and
Java

Jim Baker,
Rackspace

- Makes `sys.path` even more flexible
- site-package packages can be added to `sys.path`
- Implemented using `jar.pth`

Custom setuptools hooks

Clamp -
Seamless
integration of
Python and
Java

Jim Baker,
Rackspace

setup.py can be extended easily:

```
...
entry_points = {
    "distutils.commands": [
        "build_jar = clamp.build:build_jar",
        "singlejar = clamp.build:singlejar",
    ], ...
```

Note: these commands are now available for any package that depends on clamp!

Anatomy of a command

Clamp -
Seamless
integration of
Python and
Java

Jim Baker,
Rackspace

Enabling support for `python setup.py build_jar`, or any other custom command, requires the following attributes:

```
class build_jar(setuptools.Command):
    description = "create jar for clamped classes",
    user_options = [("output=", "o", "write jar"),]

    def initialize_options(self): ...

    def finalize_options(self): ...

    def run(self): ...
```

Adding new scripts in Python's bin

Clamp -
Seamless
integration of
Python and
Java

Jim Baker,
Rackspace

setup.py uses the same `entry_points` to point to custom scripts:

```
...
entry_points = {
    "distutils.commands": [
        "build_jar = clamp.build:build_jar",
        "singlejar = clamp.build:singlejar",
    ],
    "console_scripts": [
        "singlejar = clamp.build:singlejar_command",
    ]
}
```


Console script example: singlejar

Clamp -
Seamless
integration of
Python and
Java

Jim Baker,
Rackspace

Script is just a standard *main* function, using arg parser of choice:

```
def singlejar_command():
    parser = argparse.ArgumentParser(...)
    parser.add_argument("--classpath", ...)
    parser.add_argument("--runpy", ...)
    args = parser.parse_args()
    if args.classpath:
        args.classpath = args.classpath.split(":")
    else:
        args.classpath = []
    clamp.build.create_singlejar(
        args.output, args.classpath, args.runpy)
```

Metaprogramming to be done

Clamp -
Seamless
integration of
Python and
Java

Jim Baker,
Rackspace

- Import hooks - intercept the import of Java annotations so they can be used as class and function decorators.
- Rewriting Java bytecode - with ASM to support annotations/type signatures as class decorators.
- More, much more!

Clamp resources

Clamp -
Seamless
integration of
Python and
Java

Jim Baker,
Rackspace

- Clamp project: <https://github.com/jythontools/clamp>
- Example project: <https://github.com/jimbaker/clamped>
- Documentation: <http://clamp.readthedocs.org>